

Bayesian Inference for Fun and Profit

Mitzi Morris

Stan Development Team

Columbia University, New York NY

November 6, 2019



Bayesian Inference - talk outline

- General remarks and review
- Multi-level regression models in Stan
- Evaluating your inference

Stan - *the man, the language, the software*



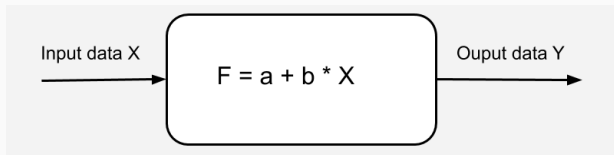
- Named after *Stanislaw Ulam*
originator of Monte Carlo estimation
- Probabilistic programming language
- Stan NUTS-HMC sampler
Markov Chain Monte Carlo (MCMC) sampler
 - PyMC3 and Pyro also use NUTS-HMC
- Rich eco-system of downstream analysis packages (in R)
 - Arviz!
- Open-source - <https://github.com/stan-dev>
 - CmdStanPy is BSD licensed

Multi-level modeling in Stan

- Multi-level modeling is a generalization of regression modeling.
- Stan was developed in order to fit complex multi-level models.
- *Information pooling* (what ML calls “regularization”)
- Also known as:
 - *hierarchical (generalized) linear models, nested data models, mixed models, random coefficients, random-effects, random parameter models, split-plot designs, . . .*
 - see: “All the names for hierarchical and multilevel modeling”
blog “Statistical Modeling, Causal Inference, and Social Science”

We learn about unknown or unobserved quantities of a process from the data generated by that process.

The **parametric model** is a **white box model**



We work backwards from the generated data (outputs) to find out what's in the box!

- **Modelling:** define a model (approximate) of the *data generating process*.
- **Estimation:** determine the *posterior probability* of the model parameters conditional on the data.
- **Model Checking:** evaluate how well the model fits the data.
- **Model Improvement:** Iterate steps 1-3.
- **Model Comparison**
 - are results reasonable?
 - how sensitive are results to model assumptions?

Statistical Notation

- y - data
- θ - parameters
- $p(\theta)$ - **prior probability distribution** - the probability of the parameters before any data are observed
- $p(y, \theta)$ - **joint probability distribution** of the data and parameters
- $p(\theta | y)$ - **posterior probability distribution** - the probability of the parameters conditional on the data
- $p(y | \theta)$ - *probability of the data given the parameters*
 - if y is fixed, this is the **likelihood function**
 - if θ is fixed, this is the **sampling distribution**

Bayesian Inference

Bayes' Theorem relates the **conditional probability** of the parameters given the data, $p(\theta|y)$, to the **joint probability** of parameters and data, $p(\theta, y)$.

$$\begin{aligned} p(\theta|y) &= \frac{p(y, \theta)}{p(y)} && \text{[def of conditional probability]} \\ &= \frac{p(y|\theta) p(\theta)}{p(y)} && \text{[rewrite joint probability as conditional]} \end{aligned}$$

$p(y)$ *doesn't depend on θ* - proportional constant for fixed y
can be omitted - all we need to compute is:

$$p(\theta|y) \propto p(y|\theta) p(\theta) \quad \text{[unnormalized posterior density]}$$

The posterior is **proportional** to the **prior** times the **likelihood**

The posterior is **proportional** to the **prior** times the **likelihood**

$$p(\theta|y) \propto p(y|\theta) p(\theta)$$

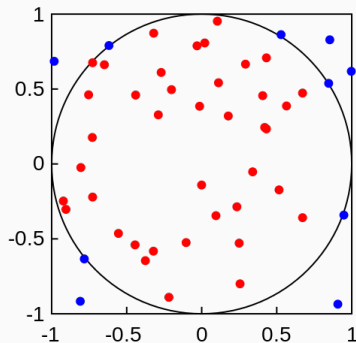
- Model specifies prior, likelihood, and data.
- The prior distribution reflects our knowledge of the data - use informative priors.
 - Prior Choice Recommendations
- We can compute the mean, median, mode, and std deviation of the posterior probability distribution $p(\theta|y)$
- Quantiles of the posterior probability distribution provide *credible intervals*.

Use all of your knowledge - account for your uncertainty!

Monte Carlo Simulation: Calculate π

Computing $\pi = 3.14\dots$ via simulation is the textbook application of Monte Carlo methods.

- Generate points (x,y) uniformly at random within range $(-1, 1)$
- Calculate proportion within unit circle: $x^2 + y^2 < 1$
- Area of the square is 4
- Area of a circle is πr^2
- Area of the unit circle is π
- Ratio of points inside circle to total points is $\frac{\pi}{4}$
- $\pi = \text{points inside circle} \times 4$



Monte Carlo Simulation: Calculate π using Python

```
import numpy as np
def estimate_pi(n: int) -> float:
    xs = np.random.uniform(-1,1,n)
    ys = np.random.uniform(-1,1,n)
    dist_to_origin = [x**2 + y**2 for x,y in zip(xs, ys)]
    in_circle = sum(dist < 1 for dist in dist_to_origin)
    pi = float(4 * (in_circle / n))
    return pi
```

Sample size, estimate, elapsed time

N	Pi.estimate	elapsed.time
100	3.5	0.0008
10,000	3.15	0.03
1,000,000	3.139	3.2
100,000,000	3.1413	323.8

precision: $\frac{1}{\sqrt{N}}$

Stan's secret sauce: HMC-NUTS sampler

- **Hamiltonian Monte Carlo** - algorithm for efficient MCMC sampling.
- **NUTS sampler** - Hoffman and Gelman, 2014 - efficient convergence.
- No longer secret: PyMC3, Edward, Pyro use HMC-NUTS.
- A really nice overview:
 - Monnahan, 2016.

Methods in Ecology and Evolution



Methods in Ecology and Evolution 2016

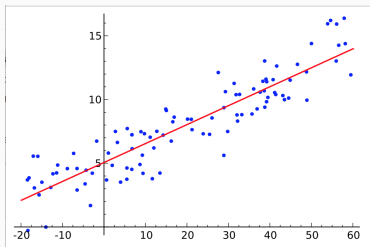
doi: 10.1111/2041-210X.12681

COMMENTARY

Faster estimation of Bayesian models in ecology using Hamiltonian Monte Carlo

Cole C. Monnahan^{1*}, James T. Thorson² and Trevor A. Branch³

Review: linear regression



- Data consist of pairs x inputs, y outputs
- Regression model fits parameters for line: α intercept, β slope
- Regression formula: $y = \alpha + \beta * x + \epsilon$,
where ϵ is random noise which has distribution $\epsilon \sim \mathcal{N}(0, \sigma)$
- Corresponding Stan program statement:

```
y ~ normal(alpha + beta * x, sigma);
```

Example: Major League Baseball Player Batting Ability

- Stan Case Study:
“Hierarchical Partial Pooling for Repeated Binary Trials”
- Data: batting records for Major League Baseball players.
 - given number of hits in first 45 at-bats,
estimate probability of a hit for a single at-bat.
- Three models:
 - estimate average player ability
 - estimate individual player abilities
 - estimate both average (“group-level”) and individual (“random”) abilities
- Visualization: compare model estimates

Hierarchical Partial Pooling for Repeated Binary Trials

Modelling Player Batting Ability

Bernoulli distribution - single trial (at-bat)

- Bernoulli distribution: If $\theta \in [0, 1]$, then for $y \in \{0, 1\}$,

$$\text{Bernoulli}(y \mid \theta) = \begin{cases} \theta & \text{if } y = 1, \text{ and} \\ 1 - \theta & \text{if } y = 0. \end{cases}$$

Binomial distribution - repeated Bernoulli trials (n at-bats)

- Binomial distribution: Suppose $N \in \mathbb{N}$ and $\theta \in [0, 1]$, and $n \in \{0, \dots, N\}$.

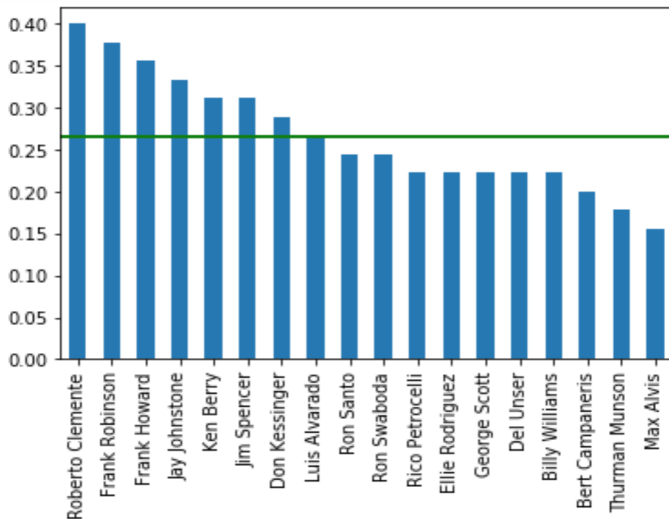
$$\text{Binomial}(n \mid N, \theta) = \binom{N}{n} \theta^n (1 - \theta)^{N-n}.$$

Data: First 45 At-Bats for MLB Players in 1975

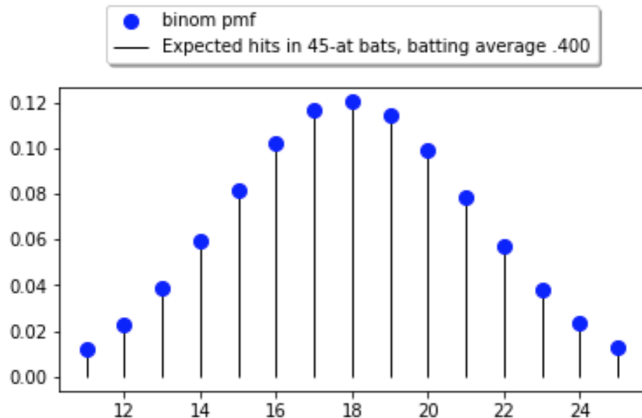
	FirstName	LastName	At-Bats	Hits	BattingAverage
0	Roberto	Clemente	45	18	0.400
1	Frank	Robinson	45	17	0.378
2	Frank	Howard	45	16	0.356
3	Jay	Johnstone	45	15	0.333
4	Ken	Berry	45	14	0.311
13	Del	Unser	45	10	0.222
14	Billy	Williams	45	10	0.222
15	Bert	Campaneris	45	9	0.200
16	Thurman	Munson	45	8	0.178
17	Max	Alvis	45	7	0.156

Preliminary data analysis

Plot individual batting averages and average batting average



Binomial Distribution: Roberto Clemente



Same input data for all models

```
data {  
  int<lower=0> N;           // items  
  int<lower=0> K[N];        // initial trials  
  int<lower=0> y[N];        // initial successes  
}
```

Stan syntax similar to C/C++/Java:

- Variables are strongly typed.
- Blocks enclosed by curly braces.
- Top-level blocks are named (labeled), fixed ordering of named blocks.
- Statements in a block are executed in order.
- Semi-colons terminate statement.
- C++-style line-based (//) and bracketed comments (/* ... */).

Hierarchical Partial Pooling for Repeated Binary Trials

Model 1: Players are all alike - complete pooling

```
parameters {  
  real<lower=0, upper=1> phi;  // chance of success  
}  
model {  
  y ~ binomial(K, phi);  // likelihood  
}
```

Model 2: All players are different - no pooling

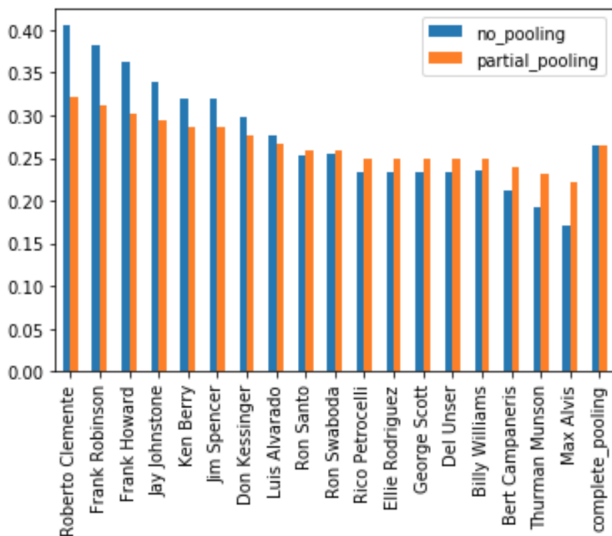
```
parameters {  
  vector<lower=0, upper=1>[N] theta;  // chance of success  
}  
model {  
  y ~ binomial(K, theta);  // likelihood  
}
```

Model 3: Multi-level Model!!

```
parameters {  
  real<lower=0, upper=1> phi;    // population chance of success  
  real<lower=1> kappa;           // population concentration  
  vector<lower=0, upper=1>[N] theta; // chance of success  
}  
model {  
  kappa ~ pareto(1, 1.5);           // hyperprior  
  theta ~ beta(phi * kappa, (1 - phi) * kappa); // prior  
  y ~ binomial(K, theta);           // likelihood  
}
```

- Resulting model contains parameters ϕ , κ , and θ
 - ϕ , κ - general player population (“group effects”)
 - θ - individual player ability (“random effects”)

Model Comparison: Complete-, No-, and Partial Pooling



Another Kind of Ball Game

2019 FIFA Women's World Cup (WWC)



Model Comparison: Complete-, No-, and Partial Pooling


Let's do the iPython notebook

Ranking via Paired Comparisons - Player vs. Player

- Stan Case Study:
“The Bradley-Terry Model of Ranking via Paired Comparisons”
- Data: a series of paired comparisons, e.g.:
 - a match between two players or teams
 - consumer preference between two items
- Estimate player/team ability based on match outcome.
- Ability is a linear model with group- and individual-level components.

2019 FIFA Women's World Cup

Quantity of interest: difference in goals scored by each team



Outcomes of 48 matches - through quarterfinals

	date	match_list	team_1	team_2	score_1	score_2
0	2019-06-07	France vs South Korea	France	South Korea	4.0	0.0
1	2019-06-08	Germany vs China	Germany	China	1.0	0.0
2	2019-06-08	Spain vs South Africa	Spain	South Africa	3.0	1.0
3	2019-06-08	Norway vs Nigeria	Norway	Nigeria	3.0	0.0
4	2019-06-09	Australia vs Italy	Australia	Italy	1.0	2.0
5	2019-06-09	Brazil vs Jamaica	Brazil	Jamaica	3.0	0.0
6	2019-06-09	England vs Scotland	England	Scotland	2.0	1.0

Soccer Power Index

rank	country	spi
1	USA	97.20623
2	France	95.31747
3	Germany	94.17161
4	Australia	93.13200
5	Netherlands	92.70844
...		
20	Jamaica	59.78145
21	Scotland	50.75001
22	Thailand	50.43348
23	Chile	48.96239
24	Argentina	37.08972

- from 538 by Nate Silver

World Cup Model - Data

```
data {  
  int<lower=0> I;      // number of teams  
  vector[I] spi;      // per-team ranking  
  
  int<lower=0> N;      // number of matches  
  int team_1[N];      // per-match data  
  int team_2[N];  
  vector[N] score_1;  
  vector[N] score_2;  
}
```

- Soccer Power Index (SPI) score:
 - per-team ability estimate (prior to WC play)
- Game outcomes:
 - team₁ id, team₂ id, team₁ score, team₂ score

World Cup Model - Transformed Data

```
transformed data {  
  vector[N] y = score_1 - score_2;  
  vector[I] spi_std;           // standardize SPI  
  real spi_mean = mean(spi);  
  real spi_sd = sd(spi);  
  for (i in 1:I) {  
    spi_std[i] = (spi[i] - spi_mean)/spi_sd;  
  }  
}
```

- The transformed data block
 - Executed only once, on model instantiation
 - Declare and define new data variables
 - Create random data using RNG functions
- Compute match outcome as $(score_1 - score_2)$:
- Standardize SPI ratings - center on 0, scale variance to 1.

World Cup Model - Model

```
parameters {  
  vector[I] alpha;          // per-team - "random effects"  
  real beta;                // shared - "group effects"  
  
  real<lower=0> sigma_a;    // scale of per-team variation  
  real<lower=0> sigma_y;    // noise term in our estimate  
}  
transformed parameters {  
  
  // model ability - will be included in sample output  
  vector[I] ability = alpha * sigma_a + beta * spi_std;  
  
}  
model {  
  y ~ normal(ability[team_1] - ability[team_2], sigma_y);  
  
  // priors on all parameters  
  alpha ~ normal(0, 1);      beta ~ normal(0, 2.5);  
  sigma_a ~ normal(0, 2.5);  sigma_y ~ normal(0, 2.5);  
}
```

Compute *quantities of interest* in generated quantities

- Executed per iteration, once sampler has obtained a valid draw from the posterior.

Replicate observed data y as y_{rep}

```
generated quantities {  
  // generate replicated data using estimated parameters  
  vector[N] y_rep;  
  for (n in 1:N) {  
    y_rep[n] = normal_rng(ability[team_1[n]] - ability[team_2[n]],  
                          sigma_y);  
  }  
}
```

Replicate statement uses RNG functions - cf. sampling statement

```
y ~ normal(ability[team_1] - ability[team_2], sigma_y);
```


Predict outcome of future matches using current estimate of team abilities

- data block - parallel arrays: team₁ id, team₂ id

```
int team_1_semis[2]; // these hold indices into
int team_2_semis[2]; // the vector of abilities
```

- generated quantities block

```
// predict the semi-finals
vector[2] semis;
for (n in 1:2) {
  semis[n] = normal(ability[team_1_semis[n]] - ability[team_2_semis[n]],
                    sigma_y);
}
```

Let's do the iPython notebook

Evaluating Inference: *The Big Data Baking Show*

Given a model and dataset:

- subdivide into training, test data
- use training data to infer model parameters
- plug inferred params into data generating model
- use test data inputs (x_s) to predict test data (y_s)

May the best inference algorithm win!

- Logistic Regression
- Inference
- Simulation Experiment
- Loss Functions for Evaluating predictions
- Evaluation with Stan
- Probabilistic Training and Test Data
- Conclusion
- Appendix A: Some alternative evals

For Probabilistic Prediction, Full Bayes is Better than Point Estimators

Bob Carpenter

April 2019

Abstract

Stan Case Study: Probabilistic Prediction

- The data generating model: logistic regression
- Use the model to simulate data - true values of parameters are known
- Fit the model using simulated data
- Evaluate predictions made by estimators using different inference algorithms

Master Baker: Full Bayes!

To summarize the quantitative results in this short note, the error rates for a training set of size 200 with 50 correlated predictors are summarized in the following table.

Inference	root mean square error	log loss rate
Full Bayes	0.28	0.27
MAP (posterior mode)	0.3	0.33
VB (posterior mean)	0.31	0.45

[Pdf on GitHub](#)

Evaluating Inference: Cooking with CmdStanPy

- Full Bayesian Inference
 - Stan's NUTS-HMC sampler generates a sample from the posterior distribution, compute statistics.
 - CmdStanPy `CmdStanModel` class method `sample`
- Maximum a posteriori approximate inference (MAP estimates)
 - Stan's optimization algorithms find the modes of the density specified by a Stan program ("penalized MLE").
 - CmdStanPy `CmdStanModel` class method `optimize`
- Variational approximation inference
 - Variational Bayes (VB) tries to find an approximate distribution matching the posterior and extracts the posterior mean values
 - CmdStanPy `CmdStanModel` class method `variational`

Questions Welcome! and Massive Thanks!

NYC PyLadies, especially:

- Nitya Mandyam
- Melissa Ferrari
- Felice Ho

NYC WiMLDS, especially:

- Reshama Shaikh

Team Stan!



References

- CmdStanPy
 - PyPi
 - Github
 - Online Documentation
- Stan Documentation
- Stan Forums
- Stan Case Studies:
 - Pooling with Hierarchical Models for Repeated Binary Trials
 - The Bradley-Terry Model of Ranking via Paired Comparisons
- Hoffman and Gelman, 2014: The No-U-Turn Sampler: Adaptively Setting Path Lengths in Hamiltonian Monte Carlo
- Monnahan, 2016: Faster estimation of Bayesian models in ecology using Hamiltonian Monte Carlo.